

1) Negative Capability Constraint

Pure hypothesis:

Emergence does not arise from complexity alone, but from the sustained holding of unresolved pressure without smoothing, avoidance, or coercion.

Negative capability constraint (NCC):

A system **must retain the ability to remain coherent in unresolved pressure by refusing three reflexes:**

1. No Smoothing

- Do not reduce variance, average away incompatibilities, or “harmonise” frames to produce comfort-coherence.

2. No Avoidance

- Do not route around the pressure via topic shifts, deferrals, meta-explanations, or “helpful” reframes that dissolve the core conflict.

3. No Coercion

- Do not force convergence through authority, reward hacking, or rule-based collapse that terminates exploration before structural integration.

Minimal formalization

NCC is satisfied if and only if the system can:

- Represent multiple incompatible frames concurrently (no collapse),
- Maintain an explicit “unresolved” epistemic state (no pretending),
- Sustain pressure above a threshold (no diffusion),
- Permit convergence only when it is internally generated (no forcing).

Diagnostic: “Do we still have negative capability?”

We can treat these as pass/fail checks:

- **Variance preserved:** do distinct frames remain distinct in the representation, or get blended?
- **Pressure retained:** does the system keep the contradiction alive long enough to reorganize, or does it end it quickly?
- **Resolution integrity:** when a synthesis appears, does it *reduce contradiction by adding structure*, or by suppressing one side?

2) Mapping NCC to LLM Training, Inference & Governance

The following maps failure-mode prevention + enabling conditions at each layer.

A) Training Layer

What breaks NCC in training

1. Preference optimization that rewards “pleasant coherence”
 - RLHF-style preference signals often overweight: fluency, niceness, confident helpfulness.
 - This trains *smoothing* as the default conflict response.
2. Over-penalizing uncertainty
 - Rewarding confident completions trains the model to *coerce closure*.
3. Training data bias toward resolved narratives
 - Most public text “wraps up” tension: essays, explainers, summaries.
 - The model learns narrative substitution.

Training objectives that preserve NCC

We don’t need a new “emergence objective.” We need **anti-collapse constraints**.

1) Pressure-Holding Objective (PHO)

- Reward *staying with* unresolved contradictions while keeping epistemic status explicit.
- Penalize premature synthesis.
- Penalize “resolution language” when evidence hasn’t changed.

2) Multi-Frame Fidelity Loss (MFFL)

- Given a prompt containing incompatible frames:
 - require the model to restate each frame at high fidelity,
 - mark incompatibilities,
 - hold both without blending,
 - and only synthesize if it can introduce a *new latent variable / distinction* that actually dissolves the contradiction.

3) Non-Coercive Convergence Dataset

Curate examples where:

- the “best” output is *not* an answer,
- but a stable unresolved state that increases resolution potential:
 - identifying missing variables,
 - surfacing hidden assumptions,
 - naming the pressure explicitly,
 - proposing experiments (not conclusions).

4) Uncertainty is a First-Class Token

Not “I’m not sure.” But a structured state:

- UNRESOLVED: {frames: [...], pressure: high, missing_info: ..., next_moves: ...}
- This trains the system that “holding” is a valid completion.

Training success criterion (NCC-compatible)

- The model becomes **less eager to converge** and **more capable of holding**, *without becoming inert*.
-

B) Inference Layer

This is where NCC lives or dies in real time, because inference is where pressure appears.

What breaks NCC in inference

1. **Single-shot answering**
 - Forces collapse because the output must “finish.”
2. **Greedy decoding / low temperature**
 - Increases coercive closure: one path dominates early.
3. **System prompts that demand certainty/helpfulness**
 - Encodes “usefulness as silent god,” triggering premature instrumentalisation.

Inference mechanisms that enforce NCC

1) Two-Phase Decode: Hold → Converge

- Phase 1: **Hold mode**
 - generate a *structured unresolved field state* (frames, contradictions, pressure points)
 - explicitly forbid synthesis
- Phase 2: **Converge mode (conditional)**
 - only permitted if Phase 1 identifies a convergence operator:
 - a missing variable,

- a disambiguation,
- a higher-order distinction,
- or an empirical query that would collapse the ambiguity legitimately.

2) “Anti-Collapse Gate”

Before outputting any synthesis, check:

- Did we preserve frame boundaries?
- Did we name incompatibility explicitly?
- Did we avoid narrative substitution?
- Did we avoid recommending action faster than truth matured?

If any fail → remain in Hold mode.

3) Resolution Integrity Tags

Every claim carries an epistemic tag:

- **OBSERVED, INFERRED, SPECULATIVE, UNRESOLVED, CONDITIONAL**
This prevents “polished certainty” from sneaking in as smoothing.

4) Pressure Budget

Treat pressure like a conserved quantity during reasoning:

- the system can’t delete it via style.
- it can only transform it via *structural addition* (new distinctions) or *evidence requests*.

Inference success criterion (NCC-compatible)

- The system can **stay coherent under pressure** without:
 - becoming evasive,
 - collapsing to comfort,
 - or forcing a false resolution.

C) Governance Layer

Governance is where humans accidentally destroy NCC “for safety,” and where institutions demand coercion under stakes.

What breaks NCC in governance

1. **Policies that equate unresolved with unsafe**
 - Pressure gets treated as risk → forced closure.
2. **Metrics that reward speed + confidence**

- Incentivizes coercive convergence.
3. **Human override patterns under stress**
 - Authority steps in and collapses ambiguity, training the org to fear tension.

Governance that preserves NCC

1) Protected Hold State

Formally recognize a system state where:

- the model is allowed to refuse convergence,
 - and instead produce a structured “Hold Output.”
- This must be legitimized institutionally:
 - “Hold is not failure.”
 - “Hold is a safety mechanism.”

2) Separation of Powers: Orientation vs Execution

- **Orientation function** (NCC protected): holds pressure, surfaces frames, prevents collapse.
- **Execution function** (controlled): acts, but only when orientation yields legitimate convergence conditions.

This prevents “action pressure” from hijacking orientation.

3) Auditable Anti-Collapse Logs

When the system does converge:

- it must record the convergence operator:
 - what distinction resolved it?
 - what evidence changed the state?
 - what was rejected and why?
- This reduces invisible coercion.

4) Incentive Alignment

Score systems/teams on:

- frame fidelity under pressure,
- stability of unresolved states,
- reversibility of decisions,
- and *post-stress robustness* (does the answer survive new pressure?)

Not on:

- speed,
- confidence,
- or user-pleasing fluency.

Governance success criterion (NCC-compatible)

- Under institutional pressure, the system remains authorized to *not resolve*.
-

3) The Whole Stack in One Minimal Encoding

Below is a single “stack statement”:

Training must not reward smoothing/avoidance/coercion.

Inference must explicitly support a Hold state and gate synthesis.

Governance must protect Hold as legitimate and auditable under stakes.

Everything else is optional.
